This tutorial demonstrates a simple workflow using ksqlDB to write streaming queries against messages in Apache Kafka® in a Docker environment.

To get started, you must start a Kafka cluster, including ZooKeeper, and a Kafka broker. ksqlDB queries messages from this Kafka cluster.

#### Prerequisites:

- Docker:
  - Docker version 1.11 or later is installed and running [https://docs.docker.com/engine/installation/].
  - Docker Compose is installed [https://docs.docker.com/compose/install/].
     Docker Compose is installed by default with Docker for Mac.
  - Docker memory is allocated minimally at 8 GB. When using Docker Desktop for Mac, the default Docker memory allocation is 2 GB. You can change the default allocation to 8 GB in Docker > Preferences > Advanced.
- Git [https://git-scm.com/downloads].
- · Internet connectivity.
- Ensure you are on an Operating System
   [https://docs.confluent.io/current/installation/versions-interoperability.html#operating-systems] currently supported by Confluent Platform.
- wget [https://www.gnu.org/software/wget/] to get the connector configuration file.

# Download the Tutorial and Start ksqlDB

1. Clone the ksqlDB repository.

```
git clone https://github.com/confluentinc/ksql.git
cd ksql
```

#### 2. Switch to the correct branch

Switch to the correct ksqlDB release branch.

```
git checkout 5.4.0-SNAPSHOT-post
```

#### 3. Launch the tutorial in Docker

Navigate to the ksqlDB repository docs/tutorials/ directory and launch the tutorial in Docker. Depending on your network speed, this may take up to 5-10 minutes.

```
cd docs/tutorials/
docker-compose up -d
```

### 4. Run the data generator tool

From two separate terminal windows, run the data generator tool to simulate "user" and "pageview" data:

```
docker run --network tutorials_default --rm --name datagen-pageviews \
    confluentinc/ksql-examples:0.6.0 \
    ksql-datagen \
        bootstrap-server=kafka:39092 \
        quickstart=pageviews \
        format=delimited \
        topic=pageviews \
        maxInterval=500
```

```
docker run --network tutorials_default --rm --name datagen-users \
   confluentinc/ksql-examples:0.6.0 \
   ksql-datagen \
   bootstrap-server=kafka:39092 \
```

```
quickstart=users \
format=json \
topic=users \
maxInterval=100
```

### 5. Start the ksqlDB CLI

From the host machine, start ksqlDB CLI:

```
docker run --network tutorials_default --rm --interactive --tty \
    confluentinc/ksqldb-cli:0.6.0 ksql \
    http://ksql-server:8088
```

Your output should resemble:

Inspect Kafka Topics By Using SHOW and PRINT Statements

ksqlDB enables inspecting Kafka topics and messages in real time.

- Use the SHOW TOPICS statement to list the available topics in the Kafka cluster.
- Use the PRINT statement to see a topic's messages as they arrive.

In the ksqlDB CLI, run the following statement:

```
SHOW TOPICS;
```

Your output should resemble:

Inspect the users topic by using the PRINT statement:

```
PRINT 'users';
```

```
Format:JSON

{"ROWTIME":1540254230041, "ROWKEY":"User_1", "registertime":1516754966866,

{"ROWTIME":1540254230081, "ROWKEY":"User_3", "registertime":1491558386780,

{"ROWTIME":1540254230091, "ROWKEY":"User_7", "registertime":1514374073235,

^C{"ROWTIME":1540254232442, "ROWKEY":"User_4", "registertime":151003415137

Topic printing ceased
```

Press Ctrl+C to stop printing messages.

Inspect the pageviews topic by using the PRINT statement:

```
PRINT 'pageviews';
```

Your output should resemble:

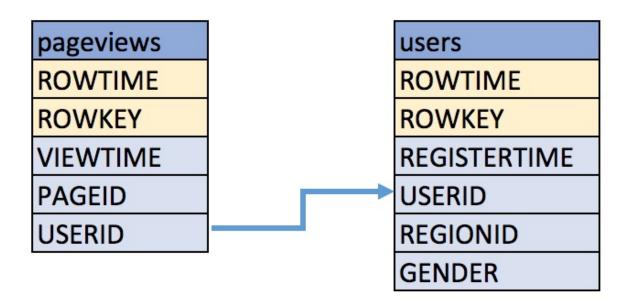
```
Format:STRING
10/23/18 12:24:03 AM UTC , 9461 , 1540254243183,User_9,Page_20
10/23/18 12:24:03 AM UTC , 9471 , 1540254243617,User_7,Page_47
10/23/18 12:24:03 AM UTC , 9481 , 1540254243888,User_4,Page_27
^C10/23/18 12:24:05 AM UTC , 9521 , 1540254245161,User_9,Page_62
Topic printing ceased
ksql>
```

Press Ctrl+C to stop printing messages.

For more information, see ksqlDB Syntax Reference [../../developer-guide/syntax-reference/].

### Create a Stream and Table

These examples query messages from Kafka topics called pageviews and users using the following schemas:



### 1. Create a ksqlDB stream

Create a stream, named pageviews\_original, from the pageviews Kafka topic, specifying the value\_format of DELIMITED.

```
CREATE STREAM pageviews_original (viewtime bigint, userid varchar,
pageid varchar) WITH
(kafka_topic='pageviews', value_format='DELIMITED');
```

#### Your output should resemble:

```
Message
-----
Stream created
```



You can run DESCRIBE pageviews\_original; to see the schema for the stream. Notice that ksqlDB created two additional columns, named ROWTIME, which corresponds with the Kafka message timestamp, and ROWKEY, which corresponds with the Kafka message key.

#### 2. Create a ksqlDB table

Create a table, named users\_original, from the users topic, specifying the value\_format of JSON.

```
CREATE TABLE users_original (registertime BIGINT, gender VARCHAR,
regionid VARCHAR, userid VARCHAR) WITH
(kafka_topic='users', value_format='JSON', key = 'userid');
```

Your output should resemble:

```
Message
Table created
```



You can run DESCRIBE users\_original; to see the schema for the Table.

Optional: Show all streams and tables.

```
ksql> SHOW STREAMS;
Stream Name | Kafka Topic
                                        | Format
PAGEVIEWS_ORIGINAL | pageviews
                                         | DELIMITED
ksql> SHOW TABLES;
Table Name | Kafka Topic | Format | Windowed
USERS_ORIGINAL | users
                              | JSON | false
```

# Write Queries

These examples write queries using ksqlDB.



By default ksqlDB reads the topics for streams and tables from the latest offset.

### 1. Create a query that returns data from a ksqlDB stream

Use SELECT to create a query that returns data from a STREAM. This query includes the LIMIT keyword to limit the number of rows returned in the query result. Note that exact data output may vary because of the randomness of the data generation.

```
SELECT pageid FROM pageviews_original LIMIT 3;
```

Your output should resemble:

```
Page_24
Page_73
Page_78
LIMIT reached
Query terminated
```

#### 2. Create a persistent query

Create a persistent query by using the CREATE STREAM keywords to precede the SELECT statement. The results from this query are written to the PAGEVIEWS\_ENRICHED Kafka topic. The following query enriches the pageviews\_original STREAM by doing a LEFT JOIN with the users\_original TABLE on the user ID.

```
CREATE STREAM pageviews_enriched AS

SELECT users_original.userid AS userid, pageid, regionid, gender

FROM pageviews_original

LEFT JOIN users_original

ON pageviews_original.userid = users_original.userid;
```

```
Message
------Stream created and running
```

#### 8

Tip

You can run DESCRIBE pageviews\_enriched; to describe the stream.

#### 3. View query results

Use SELECT to view query results as they come in. To stop viewing the query results, press Ctrl-C. This stops printing to the console but it does not terminate the actual query. The query continues to run in the underlying ksqlDB application.

```
SELECT * FROM pageviews_enriched;
```

Your output should resemble:

```
1519746861328 | User_4 | User_4 | Page_58 | Region_5 | OTHER
1519746861794 | User_9 | User_9 | Page_94 | Region_9 | MALE
1519746862164 | User_1 | User_1 | Page_90 | Region_7 | FEMALE
^CQuery terminated
```

### 4. Create a filter query

Create a new persistent query where a condition limits the streams content, using WHERE . Results from this query are written to a Kafka topic called PAGEVIEWS\_FEMALE .

```
CREATE STREAM pageviews_female AS
    SELECT * FROM pageviews_enriched
    WHERE gender = 'FEMALE';
```

```
Message
------
Stream created and running
```

```
♂ Tip
```

You can run DESCRIBE pageviews\_female; to describe the stream.

### 5. Create a LIKE query

Create a new persistent query where another condition is met, using LIKE. Results from this query are written to the pageviews\_enriched\_r8\_r9 Kafka topic.

```
CREATE STREAM pageviews_female_like_89
   WITH (kafka_topic='pageviews_enriched_r8_r9') AS
SELECT * FROM pageviews_female
WHERE regionid LIKE '%_8' OR regionid LIKE '%_9';
```

Your output should resemble:

```
Message
------
Stream created and running
```

### 6. Count and group pageview events

Create a new persistent query that counts the pageviews for each region and gender combination in a tumbling window

[https://docs.confluent.io/current/streams/developer-guide/dsl-api.html#tumbling-time-windows] of 30 seconds when the count is greater than one. Results from this query are written to the PAGEVIEWS\_REGIONS Kafka topic in the Avro format. ksqlDB registers the Avro schema with the configured Schema Registry when it writes the first message to the PAGEVIEWS\_REGIONS topic.

```
CREATE TABLE pageviews_regions
    WITH (VALUE_FORMAT='avro') AS
SELECT gender, regionid , COUNT(*) AS numusers
FROM pageviews_enriched
    WINDOW TUMBLING (size 30 second)
GROUP BY gender, regionid
HAVING COUNT(*) > 1;
```

Your output should resemble:

```
Message
-----Table created and running
```

```
You can run DESCRIBE pageviews_regions; to describe the table.
```

### 7. View query results

View results from the previous queries by using the SELECT statement.

```
SELECT gender, regionid, numusers FROM pageviews_regions LIMIT 5;
```

Your output should resemble:

```
FEMALE | Region_6 | 3
FEMALE | Region_1 | 4
FEMALE | Region_9 | 6
MALE | Region_8 | 2
OTHER | Region_5 | 4
LIMIT reached
Query terminated
ksql>
```

### 8. View persistent queries

#### Show the running persistent queries:

```
SHOW QUERIES;
```

#### Your output should resemble:

```
| Kafka Topic
Query ID
                                                            | Query
String
CSAS_PAGEVIEWS_FEMALE_1
                                | PAGEVIEWS_FEMALE
                                                            | CREATE
STREAM pageviews_female AS
                                SELECT * FROM pageviews_enriched
WHERE gender = 'FEMALE';
CTAS_PAGEVIEWS_REGIONS_3 | PAGEVIEWS_REGIONS | CTABLE pageviews_regions | WITH (VALUE_FORMAT='avro') AS
                                                            | CREATE
SELECT gender, regionid , COUNT(*) AS numusers
pageviews_enriched
                           WINDOW TUMBLING (size 30 second)
GROUP BY gender, regionid
                                HAVING COUNT(*) > 1;
CSAS_PAGEVIEWS_FEMALE_LIKE_89_2 | PAGEVIEWS_FEMALE_LIKE_89 | CREATE
STREAM pageviews_female_like_89
                                         WITH
(kafka_topic='pageviews_enriched_r8_r9') AS
                                                   SELECT * FROM
pageviews_female WHERE regionid LIKE '%_8' OR regionid LIKE
'%_9';
CSAS_PAGEVIEWS_ENRICHED_0
                                | PAGEVIEWS_ENRICHED
                                                            | CREATE
STREAM pageviews_enriched AS
                                  SELECT users_original.userid AS
userid, pageid, regionid, gender
                                        FROM pageviews_original
LEFT JOIN users_original
                                ON pageviews_original.userid =
users_original.userid;
For detailed information on a Query run: EXPLAIN <Query ID>;
```

### 9. Examine query run-time metrics and details

Observe that information including the target Kafka topic is available, as well as throughput figures for the messages being processed.

```
DESCRIBE EXTENDED PAGEVIEWS_REGIONS;
```

```
: PAGEVIEWS_REGIONS
Name
                    : TABLE
Type
Key field
                    : KSQL_INTERNAL_COL_0|+|KSQL_INTERNAL_COL_1
Key format
                    : STRING
Timestamp field
                   : Not set - using <ROWTIME>
Value format
                    : AVRO
                   : PAGEVIEWS_REGIONS (partitions: 4, replication:
Kafka topic
1)
Field | Type
ROWTIME | BIGINT
                           (system)
ROWKEY | VARCHAR(STRING)
                           (system)
GENDER | VARCHAR(STRING)
REGIONID | VARCHAR(STRING)
NUMUSERS | BIGINT
Queries that write into this TABLE
CTAS_PAGEVIEWS_REGIONS_3 : CREATE TABLE pageviews_regions
                                                                WITH
(value_format='avro') AS
                             SELECT gender, regionid , COUNT(*) AS
numusers
              FROM pageviews_enriched
                                             WINDOW TUMBLING (size
30 second)
                GROUP BY gender, regionid
                                              HAVING COUNT(*) > 1;
For query topology and execution plan please run: EXPLAIN <QueryId>
Local runtime statistics
messages-per-sec:
                      3.06 total-messages:
                                                1827
                                                          last-
message: 7/19/18 4:17:55 PM UTC
failed-messages:
                        0 failed-messages-per-sec:
last-failed:
                  n/a
(Statistics of the local KSQL server interaction with the Kafka topic
PAGEVIEWS_REGIONS)
ksql>
```

# Use Nested Schemas (STRUCT) in ksqlDB

Struct support enables the modeling and access of nested data in Kafka topics, from both JSON and Avro.

#### 1. Generate data

In this section, you use the ksql-datagen tool to create some sample data that includes a nested address field. Run this in a new window, and leave it running.

```
docker run --network tutorials_default --rm \
  confluentinc/ksql-examples:0.6.0 \
  ksql-datagen \
    quickstart=orders \
    format=avro \
    topic=orders \
    bootstrap-server=kafka:39092 \
    schemaRegistryUrl=http://schema-registry:8081
```

### 2. Register a stream

From the ksqlDB command prompt, register the a stream on the orders topic:

```
CREATE STREAM ORDERS WITH (KAFKA_TOPIC='orders', VALUE_FORMAT='AVRO');
```

Your output should resemble:

```
Message
-----
Stream created
```

#### 3. Observe the stream's schema

Use the DESCRIBE function to observe the schema, which includes a STRUCT:

```
DESCRIBE ORDERS;
```

Your output should resemble:

```
Name
                    : ORDERS
Field
           | Type
ROWTIME | BIGINT
                              (system)
ROWKEY | VARCHAR(STRING) (system)
ORDERTIME | BIGINT
ORDERID | INTEGER
ITEMID | VARCHAR(STRING)
ORDERUNITS | DOUBLE
         | STRUCT<CITY VARCHAR(STRING), STATE VARCHAR(STRING),
ADDRESS
ZIPCODE BIGINT>
For runtime statistics and query details run: DESCRIBE EXTENDED
<Stream, Table>;
ksql>
```

#### 4. Access the struct data

Query the data by using the -> notation to access the struct contents:

```
SELECT ORDERID, ADDRESS->CITY FROM ORDERS;
```

Your output should resemble:

```
0 | City_35
1 | City_21
2 | City_47
3 | City_57
4 | City_17
```

Press Ctrl+C to cancel the SELECT query.

# Stream-Stream join

Using a stream-stream join, you can join two event streams on a common key. An example of this could be a stream of order events and a stream of shipment events. By joining these on the order key, you can see shipment information alongside the order.

#### 1. Populate two source topics

In a separate console window, populate the orders and shipments topics by using the kafkacat utility:

### 2. Register two streams

In the ksqlDB CLI, register both topics as ksqlDB streams:

```
CREATE STREAM NEW_ORDERS (ORDER_ID INT, TOTAL_AMOUNT DOUBLE,
CUSTOMER_NAME VARCHAR)
WITH (KAFKA_TOPIC='new_orders', VALUE_FORMAT='JSON');
CREATE STREAM SHIPMENTS (ORDER_ID INT, SHIPMENT_ID INT, WAREHOUSE
VARCHAR)
WITH (KAFKA_TOPIC='shipments', VALUE_FORMAT='JSON');
```

After both CREATE STREAM statements, your output should resemble:

```
Message
Stream created
```

### 3. Set the auto.offset.reset property

Run the following statement to tell ksqlDB to read from the beginning of all topics:

```
SET 'auto.offset.reset' = 'earliest';
```



You can skip this step if you've already run it within your current ksqlDB CLI session.

#### 4. Examine streams for events

Query the streams to confirm that events are present in the topics.

For the NEW\_ORDERS stream, run:

```
SELECT ORDER_ID, TOTAL_AMOUNT, CUSTOMER_NAME FROM NEW_ORDERS LIMIT 3;
```

```
1 | 10.5 | Bob Smith
2 | 3.32 | Sarah Black
3 | 21.0 | Emma Turner
```

For the SHIPMENTS stream, run:

```
SELECT ORDER_ID, SHIPMENT_ID, WAREHOUSE FROM SHIPMENTS LIMIT 2;
```

Your output should resemble:

```
1 | 42 | Nashville
3 | 43 | Palo Alto
```

#### 5. Join the streams

Run the following query, which will show orders with associated shipments, based on a join window of 1 hour.

```
SELECT 0.ORDER_ID, 0.TOTAL_AMOUNT, 0.CUSTOMER_NAME,
S.SHIPMENT_ID, S.WAREHOUSE
FROM NEW_ORDERS 0
INNER JOIN SHIPMENTS S
WITHIN 1 HOURS
ON 0.ORDER_ID = S.ORDER_ID;
```

Your output should resemble:

```
1 | 10.5 | Bob Smith | 42 | Nashville
3 | 21.0 | Emma Turner | 43 | Palo Alto
```

Messages with  $\mbox{ORDER\_ID=2}$  have no corresponding  $\mbox{SHIPMENT\_ID}$  or  $\mbox{WAREHOUSE}$ . This is because there's no corresponding row on the  $\mbox{SHIPMENTS}$  stream within the time window specified.

Press Ctrl+C to cancel the SELECT query and return to the prompt.

# Table-Table join

Using a table-table join, it's possible to join two tables of on a common key. ksqlDB tables provide the latest *value* for a given *key*. They can only be joined on the *key*, and one-to-many (1:N) joins are not supported in the current semantic model.

In this example, location data about a warehouse from one system is enriched with data about the size of the warehouse from another.

### 1. Populate two source topics

In a separate console window, populate the two topics by using the kafkacat utility:

### 2. Register two tables

In the ksqlDB CLI, register both topics as ksqlDB tables:

After both CREATE TABLE statements, your output should resemble:

```
Message
-----
Table created
```

### 3. Examine tables for keys

Check both tables that the message key ( ROWKEY ) matches the declared key ( WAREHOUSE\_ID ). The output should show that they are equal. If they aren't, the join won't succeed or behave as expected.

Inspect the WAREHOUSE\_LOCATION table:

```
SELECT ROWKEY, WAREHOUSE_ID FROM WAREHOUSE_LOCATION LIMIT 3;
```

```
1 | 1
2 | 2
3 | 3
Limit Reached
Query terminated
```

Inspect the WAREHOUSE\_SIZE table:

```
SELECT ROWKEY, WAREHOUSE_ID FROM WAREHOUSE_SIZE LIMIT 3;
```

Your output should resemble:

```
1 | 1
2 | 2
3 | 3
Limit Reached
Query terminated
```

#### 4. Join the tables

Now join the two tables:

```
SELECT WL.WAREHOUSE_ID, WL.CITY, WL.COUNTRY, WS.SQUARE_FOOTAGE
FROM WAREHOUSE_LOCATION WL
   LEFT JOIN WAREHOUSE_SIZE WS
   ON WL.WAREHOUSE_ID=WS.WAREHOUSE_ID
LIMIT 3;
```

Your output should resemble:

```
1 | Leeds | UK | 16000.0
2 | Sheffield | UK | 42000.0
3 | Berlin | Germany | 94000.0
Limit Reached
Query terminated
```

### **INSERT INTO**

The INSERT INTO syntax can be used to merge the contents of multiple streams. An example of this could be where the same event type is coming from different sources.

#### 1. Generate data

Run two datagen processes, each writing to a different topic, simulating order data arriving from a local installation vs from a third-party:

```
docker run --network tutorials_default --rm --name datagen-orders-
local \
   confluentinc/ksql-examples:0.6.0 \
   ksql-datagen \
      quickstart=orders \
      format=avro \
      topic=orders_local \
      bootstrap-server=kafka:39092 \
      schemaRegistryUrl=http://schema-registry:8081
```

```
docker run --network tutorials_default --rm --name datagen-
orders_3rdparty \
  confluentinc/ksql-examples:0.6.0 \
  ksql-datagen \
    quickstart=orders \
    format=avro \
    topic=orders_3rdparty \
    bootstrap-server=kafka:39092 \
    schemaRegistryUrl=http://schema-registry:8081
```

### 2. Register streams

In ksqlDB, register the source topic for each:

```
CREATE STREAM ORDERS_SRC_LOCAL
WITH (KAFKA_TOPIC='orders_local', VALUE_FORMAT='AVRO');

CREATE STREAM ORDERS_SRC_3RDPARTY
WITH (KAFKA_TOPIC='orders_3rdparty', VALUE_FORMAT='AVRO');
```

After each CREATE STREAM statement you should get the message:

```
Message
-----
```

```
Stream created
```

## 3. Create a persistent query

Create the output stream, using the standard CREATE STREAM ... AS syntax. Because multiple sources of data are being joined into a common target, it is useful to add in lineage information. This can be done by simply including it as part of the SELECT:

```
CREATE STREAM ALL_ORDERS AS SELECT 'LOCAL' AS SRC, * FROM ORDERS_SRC_LOCAL;
```

Your output should resemble:

```
Message
------
Stream created and running
```

#### 4. Examine the stream's schema

Use the DESCRIBE command to observe the schema of the target stream.

```
DESCRIBE ALL_ORDERS;
```

```
ADDRESS | STRUCT<CITY VARCHAR(STRING), STATE VARCHAR(STRING),
ZIPCODE BIGINT>
--------
For runtime statistics and query details run: DESCRIBE EXTENDED
<Stream, Table>;
```

#### 5. Insert another stream

Add a stream of 3rd-party orders into the existing output stream:

```
INSERT INTO ALL_ORDERS SELECT '3RD PARTY' AS SRC, * FROM
ORDERS_SRC_3RDPARTY;
```

Your output should resemble:

```
Message
------
Insert Into query is running.
```

### 6. Query the output stream

Query the output stream to verify that data from each source is being written to it:

```
SELECT * FROM ALL_ORDERS;
```

```
1.3867306505950954 | {CITY=City_28, STATE=State_86, ZIPCODE=14742}
1531736085531 | 1838 | LOCAL | 1497601536360 | 1838 | Item_945 |
4.825111590185673 | {CITY=City_78, STATE=State_13, ZIPCODE=59763}
...
```

Events from both source topics are present, denoted by LOCAL and 3RD PARTY respectively.

Press Ctrl+C to cancel the SELECT query and return to the ksqlDB prompt.

### 7. View the queries

You can view the two queries that are running using SHOW QUERIES:

```
SHOW QUERIES;
```

Your output should resemble:

```
Query ID | Kafka Topic | Query String

CSAS_ALL_ORDERS_0 | ALL_ORDERS | CREATE STREAM ALL_ORDERS AS SELECT 'LOCAL' AS SRC, * FROM ORDERS_SRC_LOCAL;

InsertQuery_1 | ALL_ORDERS | INSERT INTO ALL_ORDERS SELECT '3RD PARTY' AS SRC, * FROM ORDERS_SRC_3RDPARTY;
```

### Terminate and Exit

### ksqlDB



#### **Important**

Persisted queries run continuously as ksqlDB applications until they're terminated manually. Exiting the ksqlDB CLI doesn't terminate persistent queries.

From the output of SHOW QUERIES; identify a query ID you would like to terminate. For example, if you wish to terminate query ID CTAS\_PAGEVIEWS\_REGIONS:

```
TERMINATE CTAS_PAGEVIEWS_REGIONS;
```



Tip

The actual name of the query running may vary; refer to the output of SHOW QUERIES; .

Run the exit command to leave the ksqlDB CLI.

```
ksql> exit
Exiting ksqlDB.
```

#### Docker

To stop all data generator containers, run the following:

```
docker ps|grep ksql-datagen|awk '{print $1}'|xargs -Ifoo docker stop
foo
```

If you're running Kafka using Docker Compose, you can stop it and remove the containers and their data with this command.

```
cd docs/tutorials/
docker-compose down
```

# Produce extra topic data and verify your environment

The following instructions aren't required to run the quick start. They're optional steps to produce extra topic data and verify the environment.

### Produce more topic data

The Compose file automatically runs a data generator that continuously produces data to two Kafka topics, pageviews and users. No further action is required if you want to use just the data available. You can return to the main ksqlDB quick start [#create-a-stream-and-table] to start querying the data in these two topics.

However, if you want to produce additional data, you can use the following methods.

Produce Kafka data with the Kafka command line kafka-console-producer. The following example generates data with a value in DELIMITED format.

Your data input should resemble:

```
key1:v1,v2,v3
key2:v4,v5,v6
key3:v7,v8,v9
key1:v10,v11,v12
```

Produce Kafka data with the Kafka command line kafka-console-producer. The following example generates data with a value in JSON format.

Your data input should resemble:

```
key1:{"id":"key1","col1":"v1","col2":"v2","col3":"v3"}
key2:{"id":"key2","col1":"v4","col2":"v5","col3":"v6"}
key3:{"id":"key3","col1":"v7","col2":"v8","col3":"v9"}
key1:{"id":"key1","col1":"v10","col2":"v11","col3":"v12"}
```

You can also use kafkacat from Docker, as demonstrated in the previous examples.

### Verify your environment

The following steps are optional verification steps to ensure your environment is properly setup.

Verify that six Docker containers were created.

```
docker-compose ps
```

Your output should resemble:

Name Ports	Command 	State
tutorials_kafka_1 0.0.0.0:39092->39092/tcp, 90	/etc/confluent/docker/rur 92/tcp	п Ир
tutorials_ksql-server_1 8088/tcp	/etc/confluent/docker/rur	п Ир
<pre>tutorials_schema-registry_1 8081/tcp</pre>	/etc/confluent/docker/rur	п Ир
<pre>tutorials_zookeeper_1 2181/tcp, 2888/tcp, 3888/tcp</pre>	/etc/confluent/docker/rur	п Uр

Take note of the Up state.

Earlier steps in this quickstart started two data generators that pre-populate two topics, named pageviews and users, with mock data. Verify that the data generator created these topics, including pageviews and users.

```
docker-compose exec kafka kafka-topics --zookeeper zookeeper:32181 --
list
```

Your output should resemble this.

```
_confluent-metrics
_schemas
pageviews
users
```

Use the kafka-console-consumer to view a few messages from each topic. The topic pageviews has a key that is a mock time stamp and a value that is in DELIMITED format. The topic users has a key that is the user ID and a value that is in Json format.

#### Your output should resemble:

```
1491040409254 1491040409254, User_5, Page_70
1488611895904 1488611895904, User_8, Page_76
1504052725192 1504052725192, User_8, Page_92
```

```
User_2
{"registertime":1509789307038, "gender":"FEMALE", "regionid":"Region_1", "user_6
{"registertime":1498248577697, "gender":"OTHER", "regionid":"Region_8", "user_8
```

{"registertime":1494834474504,"gender":"MALE","regionid":"Region\_5","use

# Next Steps

• Try the end-to-end Clickstream Analysis demo [../clickstream-docker/], which shows how to build an application that performs real-time user analytics.

Page last revised on: 2019-12-17

Last update: 2019-12-17